

claim. Not a side deck.

Some might say that Model had made up his mind to quit the industry before he even started. But he says he never thought of leaving until he had been in the industry for six years. "I don't know if I was wrong or not," he says. "I'm a design engineer. I've been doing that for sixteen years of industrial design before that. If a line here is wrong, the error is

It was special, I really saw counter-arguments, but it was the only time I had a chance to make my point. I was asked whether a conclusion on infrastructure and evidence is support COVID origins, LHC safety, or egg nutrition.

2. What the system does, in four stages

Four stages. Separate modules. Run in order.

Stage 1 Resolve identity (genealogy)

Code: src/epistemic/ingest.js, genealogy.js, source\_identity.js

Load a case folder (evidence\_blocks.json, source\_registry.json). Map every block to three

levels:

Level 1: Counts. Example quote in the JSON:

Lineage 1: One underlying observation. One experiment, one dataset, one generative event

Headline number is always lineage. Counting claims overstates by how many excerpts you

cite about. Lineage is the conservative count. Overstating independence is the failure mode we

live shape (COVID, after this build): excerpts documents lineages, with a citation

inflation factor printed on the report.

Stage 2 Propose

Code: src/epistemic/adjudicate.js (propose), llm.js

Ask a question of the corpus. A local model (Ollama) returns JSON: conclusion, confidence,

premises on.

That output is a proposal. Not a verdict.

Stage 3 Verify (deterministic)

Code: verifyCitations() in adjudicate.js, measure.js

No model in this stage. For each reasoning step:

1. Cited block IDs must exist in the corpus.  
2. Measurement blocks: re-run a whitelisted read-only command and compare declared value to measured value. Exact match  
Cite without quote flagged weak. All weak verdict unsubstantiated (not a pass).  
If a measurement can't run on this machine (repo not present), result is unverifiable\_here.  
Third state: Not the same as failed. Unchecked never counts as verified.  
Stage 4 Challenge panel + verdict  
Code: mechanical\_challenge.js, rest of adjudicate.js, model\_identity.js  
Two seats on the panel.  
Seat A: mechanical challenger. Always runs. Zero models. Eight checks (G1-G8); coverage  
and coverage can change from the model independence test as a display of the code, not a  
calculation. This is the only seat whose independence we can guarantee.  
Seat B: model challenger (optional). sees the conclusion, not the reasoning (blind). Picked  
from the proposers. Same family as second call, not second source. It is verified different  
Verdicts (computed, not declared):  
Verdict: 1. Meaningful  
unsubstantiated check failed. Done  
verified, unchecked. 2. Calculated and challenge ran  
contested. 3. Two or more independent lineages objected.  
Human is the last message on the panel. Recording a position raises independence to strong.  
Positions write to decisions/.json. On the record is real, not copy.

3. How that answers the bug

Overcounting. Stage 1 reports lineages, not citation counts. Cochrane's Handbook (Ch. 5) says: "We don't invent that problem. We shipped tooling for the fall that stays manual. Unchallenged conclusions. Stages 2-4 force shown work, mechanical citation checks, then a self-check. Case self-prints the tool at claims about me. Mechanical challenger catches overstatements in my own evidence pages, flagged. Every entry names a test: npm run verify:register fails if a claimed guard is missing."

#### 4. What it does / doesn't (honest)

Does: solve claim, document lineage and print the conservative count  
Doesn't: persist human positions  
Does: papers (Elicit)  
Doesn't: settle the underlying scientific questions  
Does: claim stays narrow: structure and checkable independence, not better prose and not domain expertise.

#### 5. What's in the repo now (verify, don't trust me)

As of this build: npm test 188 passing in the epistack repo (no network). Re-run before

**You review**

Please note Where COVER NOTE TO ELF.md  
has been added under /docs/  
Mechanical half runs cold • Secret epistemic/mechanical challenge.js'  
...pushes https://github.com/brandon255/doctrine-labs-elf-epistack.git  
npm run epistemic:covid  
No install beyond Node 20. No API key. Mechanical half runs cold. Local model optional for live propose/challenge. Prefer not to clone open docs/transcripts or GitHub.

Circles: Where does the COVER NOTE TO ELF.md  
 go? Should we have a blank file called fffrpf-REGISTEREDPART.md?  
 Mechanical checks C1C8 • Site epistemic/mechanical\_challenge.js?  
 GitHub repo: https://github.com/brandon255/doctrine-labs-flf-epistack.git

No install beyond Node 20. No API key. Mechanical half runs cold. Local model optional for live propose/challenge. Prefer not to clone: open docs/transcripts/ on GitHub.

6. What I claim / what I don't  
learned little JavaScript from the or on top. Tests pass. Failures are named.  
What do not call it. That advice abashed too is worse than C++ programs handbooks. Senior  
in fact submitting an ad on stackoverflow.com asking for a building and you can read.  
Thanks again for the Jun 30 feedback. It moved the work.

Brandon

What I claim: I found a real gap while dogfooding, built the navigation layer you reviewed, bearing little resemblance to mine, or on top. Turns tests pass. Failures are named. I'm not submitting as a poster boy, first to voice the problem (Greenberg, 2009; Goehring handbook). Senior Infrastructure Eng while doing other work, then kept building after your early read. Thanks again for the Jun 30 feedback. It moved the work.

\*Doctrine Labs • co-written with Cursor (Rivet) • Brandon Flores\*